

第三次作业

3.1

(1) 所写 C 程序如下:

```
#include <stdio.h>
#include <time.h>
#include <stdlib.h>
#include <unistd.h>

int main(){
    int array[1000000];
    for (int i = 0; i < 1000000; i++)
        array[i] = i+1;
    long sum = 0;
    struct timespec start, end;
    clock_gettime(CLOCK_MONOTONIC, &start);
    for (int i = 0; i < 1000000; i++)
        sum += array[i];
    clock_gettime(CLOCK_MONOTONIC, &end);
    printf("Sum: %ld\n", sum);
    printf("Time: %ld ns\n", (end.tv_sec - start.tv_sec) * 1000000000 +
        (end.tv_nsec - start.tv_nsec));
    return 0;
}
```

编译运行后,输出结果如下:

```
zhangjiawei@OS:/Users/zhangjiawei/Desktop/Operating_System/hw3$ cd "/Users/zhangjiawei/Desktop/Operating_System/hw3/" && gcc 3_1_1.c -o 3_1_1 && "/Users/zhangjiawei/Desktop/Operating_System/hw3/"3_1_1
Sum: 500000500000
Time: 2421223 ns
```

图 1.3_1_1 运行结果

(2) 所写 C 程序如下:

```
#include <stdio.h>
#include <time.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>
#define THREADS 4
long sum[THREADS] = {0};
int array[1000000];
```

```
typedef struct {
    int *array;
    int id;
} thread_data_t;

void* thread_sum(void* arg){
    thread_data_t *data = (thread_data_t*)arg;
    int *array = data->array;
    int id = data->id;
    long local_sum = 0;
    int start = id * 1000000 / THREADS;
    int end = (id + 1) * 1000000 / THREADS;
    for (int i = start; i < end ; i++)
        local_sum += array[i];
    sum[id] = local_sum;
    return NULL;
}

int main(){
    for (int i = 0; i < 1000000; i++)
        array[i] = i+1;

    struct timespec start, end;
    pthread_t threads[THREADS];
    thread_data_t thread_data[THREADS];
    clock_gettime(CLOCK_MONOTONIC, &start);
    for (int i = 0; i < THREADS; i++){
        thread_data[i].array = array;
        thread_data[i].id = i;
        pthread_create(&threads[i], NULL, thread_sum, &thread_data[i]);
    }
    for (int i = 0; i < THREADS; i++)
        pthread_join(threads[i], NULL);

    long total_sum = 0;
    for (int i = 0; i < THREADS; i++)
        total_sum += sum[i];
    clock_gettime(CLOCK_MONOTONIC, &end);
    printf("Sum: %ld\n", total_sum);
    printf("Time: %ld ns\n", (end.tv_sec - start.tv_sec) * 1000000000 +
        (end.tv_nsec - start.tv_nsec));
    return 0;
}
```

```
}

```

在这个程序中，我将数组分成了四份，分别由四个线程进行求和，最后将四个线程的结果相加得到最终结果。

编译运行后，输出结果如下：

```
zhangjiawei@OS:/Users/zhangjiawei/Desktop/Operating_System/hw3$ cd "/Users/zhangjiawei/Desktop/Operating_System/hw3/" && gcc 3_1_2
.c -o 3_1_2 && "/Users/zhangjiawei/Desktop/Operating_System/hw3/"3_1_2
Sum: 500000500000
Time: 5498159 ns
zhangjiawei@OS:/Users/zhangjiawei/Desktop/Operating_System/hw3$ cd "/Users/zhangjiawei/Desktop/Operating_System/hw3/" && gcc 3_1_2
.c -o 3_1_2 && "/Users/zhangjiawei/Desktop/Operating_System/hw3/"3_1_2
Sum: 500000500000
Time: 5666289 ns
zhangjiawei@OS:/Users/zhangjiawei/Desktop/Operating_System/hw3$ cd "/Users/zhangjiawei/Desktop/Operating_System/hw3/" && gcc 3_1_2
.c -o 3_1_2 && "/Users/zhangjiawei/Desktop/Operating_System/hw3/"3_1_2
Sum: 500000500000
Time: 5573495 ns
```

图 2.3_1_2 运行结果

可以看出，四线程的运行时间大于直接进行求和的运行时间，这是因为线程的创建和销毁也需要时间，而且线程之间的切换也会带来额外的开销。经过进一步试验，发现随着线程数目从 1 开始逐渐增加，程序运行时间先降低后增加，这是因为我的 CPU 是 8 核的，当线程数目增加到 8 时，线程数目超过了 CPU 的核心数，导致线程切换的开销增加，从而使得程序运行时间增加。

(3) 所写 C 程序如下：

```
#define _GNU_SOURCE
// There may be some wrong with the homework document, if the macro is
// __USE_GNU, the program will not run. After asking Github copilot, I
// change it to _GNU_SOURCE, and the program runs successfully. Others
// also have the same problem, so I hope the teacher can check out the
// problem.

#include <stdio.h>
#include <time.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>
#include <sched.h>
#define THREADS 4
int array[1000000];
long sum[THREADS] = {0};

typedef struct {
    int *array;
    int id;
} thread_data_t;
```

```
void* thread_sum(void* arg){
    thread_data_t *data = (thread_data_t*)arg;
    int *array = data->array;
    int id = data->id;
    long local_sum = 0;
    int start = id * 1000000 / THREADS;
    int end = (id + 1) * 1000000 / THREADS;
    for (int i = start; i < end; i++)
        local_sum += array[i];
    sum[id] = local_sum;
    return NULL;
}

void* worker(void* arg){
    thread_data_t *data = (thread_data_t*)arg;
    int core_id = data->id;
    cpu_set_t cpuset;
    CPU_ZERO(&cpuset);
    CPU_SET(core_id, &cpuset);
    sched_setaffinity(0, sizeof(cpuset), &cpuset);
    thread_sum(arg);
    return NULL;
}

int main(){
    for (int i = 0; i < 1000000; i++)
        array[i] = i+1;

    struct timespec start, end;
    pthread_t threads[THREADS];
    thread_data_t thread_data[THREADS];
    clock_gettime(CLOCK_MONOTONIC, &start);
    for (int i = 0; i < THREADS; i++){
        thread_data[i].array = array;
        thread_data[i].id = i;
        pthread_create(&threads[i], NULL, worker, &thread_data[i]);
    }
    for (int i = 0; i < THREADS; i++)
        pthread_join(threads[i], NULL);

    long total_sum = 0;
    for (int i = 0; i < THREADS; i++)
        total_sum += sum[i];
}
```

```

    clock_gettime(CLOCK_MONOTONIC, &end);
    printf("Sum: %ld\n", total_sum);
    printf("Time: %ld ns\n", (end.tv_sec - start.tv_sec) * 1000000000 +
        (end.tv_nsec - start.tv_nsec));
    return 0;
}

```

在这个程序中, 我将线程绑定到了特定的 CPU 核心上, 这样可以避免线程切换的开销。编译运行后, 输出结果如下:

```

zhangjiawei@05:/Users/zhangjiawei/Desktop/Operating_System/hw3$ cd "/Users/zhangjiawei/Desktop/Operating_System/hw3/" && gcc 3_1_3
.c -o 3_1_3 && "/Users/zhangjiawei/Desktop/Operating_System/hw3/"3_1_3
Sum: 500000500000
Time: 5593189 ns
zhangjiawei@05:/Users/zhangjiawei/Desktop/Operating_System/hw3$ cd "/Users/zhangjiawei/Desktop/Operating_System/hw3/" && gcc 3_1_3
.c -o 3_1_3 && "/Users/zhangjiawei/Desktop/Operating_System/hw3/"3_1_3
Sum: 500000500000
Time: 5693191 ns
zhangjiawei@05:/Users/zhangjiawei/Desktop/Operating_System/hw3$ cd "/Users/zhangjiawei/Desktop/Operating_System/hw3/" && gcc 3_1_3
.c -o 3_1_3 && "/Users/zhangjiawei/Desktop/Operating_System/hw3/"3_1_3
Sum: 500000500000
Time: 5620856 ns

```

图 3.3_1_3 运行结果

运行的结果与上一问相差无几, 这可能是由于操作系统在上一问中自动进行了线程绑定, 所以运行时间没有太大的变化。

3.2

所写 C 程序如下:

```

#include <stdio.h>
#include <time.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>
#define THREADS 4 // Number of threads
int array[1000000];
struct thread_data_t
{
    int *array;
    int id;
}; // Structure to pass data to threads

// Function to assign values to array
void* thread_assign(void* arg){
    struct thread_data_t *data = (struct thread_data_t*)arg; // Get data from
    argument
    int *array = data->array; // Get array from data
}

```

```
int id = data->id; // Get id from data to calculate start and end index of
                  // current thread
int start = id * 1000000 / THREADS; // Calculate start index
int end = (id + 1) * 1000000 / THREADS; // Calculate end index
for (int i = start; i < end; i++)
    array[i] = i+1;
return NULL;
}

int main(){
    struct thread_data_t thread_data[THREADS]; // Array of thread data to pass to
                                                // threads
    pthread_t threads[THREADS]; // Array of threads to store thread ids
    for (int i = 0; i < THREADS; i++){
        thread_data[i].array = array;
        thread_data[i].id = i;
        pthread_create(&threads[i], NULL, thread_assign, &thread_data[i]); //
        // Create threads and pass data to them
    }
    for (int i = 0; i < THREADS; i++)
        pthread_join(threads[i], NULL); // Wait for threads to finish

    long sum = 0;
    for (int i = 0; i < 1000000; i++)
        sum += array[i];
    printf("Sum: %ld\n", sum);
    return 0;
}
```

在这个程序中,我将数组分成了四份,分别由四个线程进行赋值,最后将整个数组的值相加得到最终结果。编译运行后,输出结果如下:

```
zhangjiawei@05:/Users/zhangjiawei/Desktop/Operating_System/hw3$ cd "/Users/zhangjiawei/Desktop/Operating_System/hw3/" && gcc 3_2.c
-o 3_2 && "/Users/zhangjiawei/Desktop/Operating_System/hw3/"3_2
Sum: 500000500000
```

图 4.3_2 运行结果